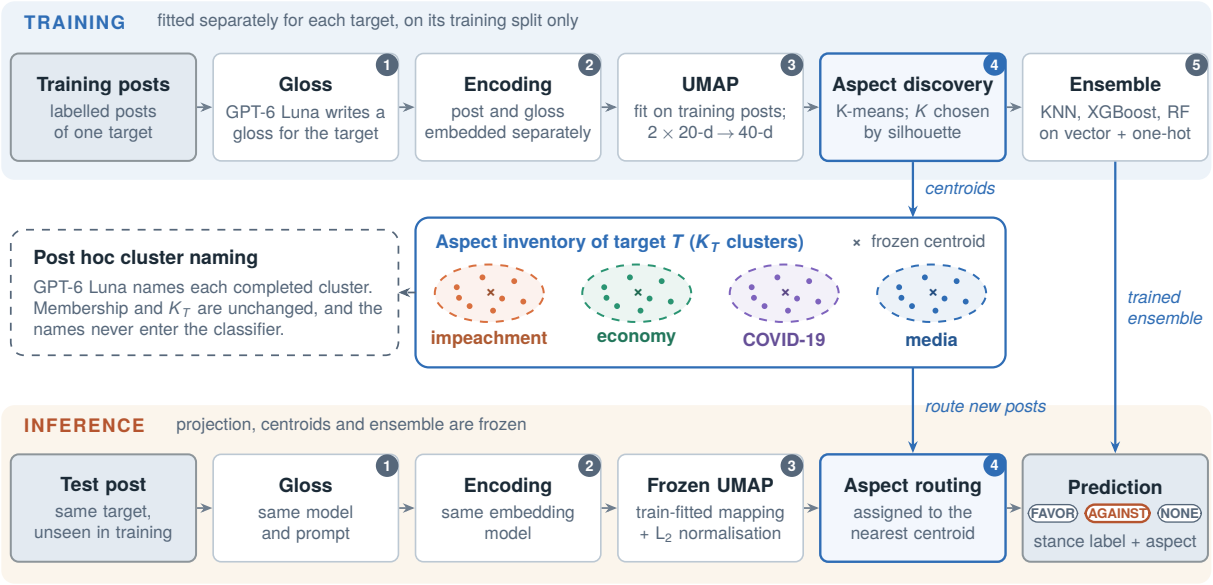
TRAINING
fitted separately for each target, on its training split only
Training posts
labelled posts of one target
Gloss
1
GPT-6 Luna writes a gloss for the target
Encoding
2
post and gloss embedded separately
UMAP
3
fit on training posts; 2 × 20-d → 40-d
Aspect discovery
4
K-means; K chosen by silhouette
Ensemble
5
KNN, XGBoost, RF on vector + one-hot
centroids
trained ensemble
Post hoc cluster naming
GPT-6 Luna names each completed cluster. Membership and K_T are unchanged, and the names never enter the classifier.
Aspect inventory of target T (K_T clusters)
× frozen centroid
impeachment
economy
COVID-19
media
route new posts
INFERENCE
projection, centroids and ensemble are frozen
Test post
same target, unseen in training
Gloss
1
same model and prompt
Encoding
2
same embedding model
Frozen UMAP
3
train-fitted mapping + L_2 normalisation
Aspect routing
4
assigned to the nearest centroid
Prediction
FAVOR
AGAINST
NONE
stance label + aspect